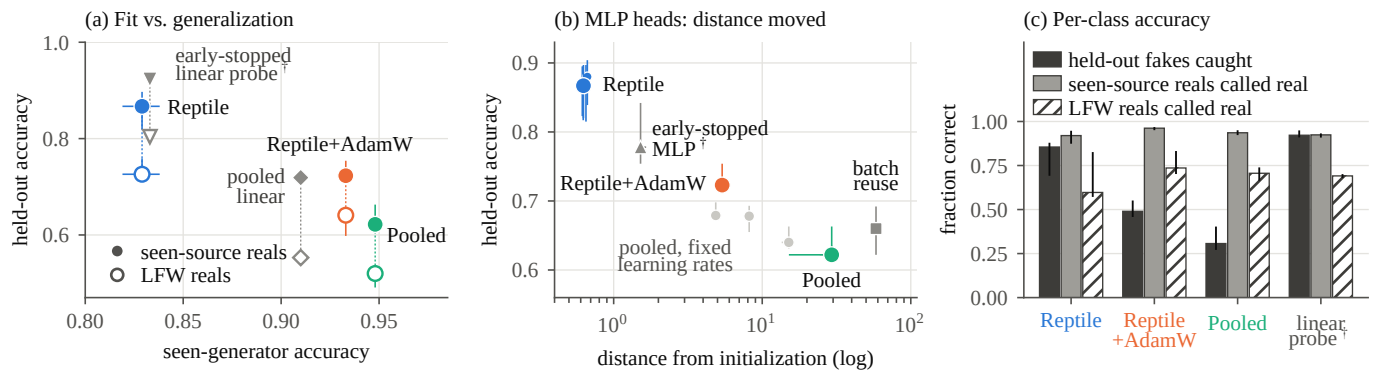
(a) Fit vs. generalization
held-out accuracy
seen-generator accuracy
1.0
0.8
0.6
0.80
0.85
0.90
0.95
early-stopped
linear probe †
Reptile
Reptile+AdamW
pooled
linear
Pooled
seen-source reals
LFW reals
(b) MLP heads: distance moved
held-out accuracy
distance from initialization (log)
0.9
0.8
0.7
0.6
10^0
10^1
10^2
Reptile
early-stopped
MLP †
Reptile+AdamW
batch
reuse
pooled, fixed
learning rates
Pooled
(c) Per-class accuracy
fraction correct
held-out fakes caught
seen-source reals called real
LFW reals called real
1.00
0.75
0.50
0.25
0.00
Reptile
Reptile
+AdamW
Pooled
linear
probe †